



# ARTIFICIAL INTELLIGENCE

---

**Dr. Rajani S**

Department of Computer Applications

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

# Problem Statement

A financial analyst wants to predict the closing price of a stock using its opening price, previous day's closing price, and trading volume.

You are provided with a historical stock dataset. Build a Linear Regression model to predict the closing price and evaluate its performance.

## Dataset

Save the following as `stock_price_prediction.csv`

| Open | Previous_Close | Volume | Close |
|------|----------------|--------|-------|
| 100  | 98             | 150000 | 104   |
| 104  | 104            | 165000 | 106   |
| 106  | 106            | 170000 | 108   |
| 108  | 108            | 180000 | 110   |

### Questions

1. Load the dataset using Pandas.
2. Display the first five records.
3. Identify the input features and target variable.
4. Split the dataset into training and testing sets.
5. Train a Linear Regression model.
6. Predict the closing prices.
7. Compare actual and predicted values.
8. Calculate MAE, MSE, RMSE and  $R^2$  Score.
9. Plot Actual vs Predicted values.

# Step 1: Import Required Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

## Explanation

Pandas → Read and manipulate the dataset.

NumPy → Numerical operations.

Matplotlib → Plot graphs.

train\_test\_split → Split data into training and testing.

LinearRegression → Build the regression model.

metrics → Evaluate the model.

## Stock\_Price\_Prediction

- stock\_prediction.py
- stock\_price\_prediction.csv

$$Y = c + m_1X_1 + m_2X_2 + m_3X_3$$

Open ( $X_1$ )

Previous\_Close ( $X_2$ )

Volume ( $X_3$ )

Close ( $Y$ )

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

Linear Regression works well only when the relationship between variables is linear (straight line).

However, many real-world datasets have nonlinear (curved) relationships.

### **Examples:**

Study Hours vs Marks

Fertilizer vs Crop Yield

Age vs Income

Sales vs Advertisement Cost

If the data forms a curve, a straight line cannot fit it accurately.

### Use Polynomial Regression when:

1. The scatter plot shows a curved pattern.
2. The relationship between variables is nonlinear.
3. Polynomial regression fits a curve instead of a straight line by adding higher-degree terms

The degree of a polynomial is the highest exponent (power) of the independent variable  $X$ .

### Degree 1 (Linear Regression)

Equation:

$$y = mX + c$$

power of  $X = 1$

Graph: Straight Line

### Degree 2 (Quadratic Polynomial)

Equation:

$$y = c + mX + aX^2$$

Graph: Curved (Parabola)

power of  $X = 2$

### Degree 3 (Cubic Polynomial)

Equation:

Highest power = 3

Graph : create more complex curves. (S shape)

$$y = c + mX + aX^2 + bX^3$$


Linear Regression uses Least Squares to fit a straight-line equation.

Polynomial Regression also uses Least Squares, but it fits a polynomial equation.

# ARTIFICIAL INTELLIGENCE


## Polynomial Regression

### 1. Linear Regression (Poor Fit for Curved Data)



✗ The straight line cannot capture the curved pattern.

### 2. Polynomial Regression (Better Fit for Curved Data)



✓ The curve follows the actual pattern of the data.

# ARTIFICIAL INTELLIGENCE

## Polynomial Regression

Dataset

| X | Y  |
|---|----|
| 1 | 6  |
| 2 | 11 |
| 3 | 18 |

Assume

$$y=c+mX+aX^2$$

### First point

$$X=1$$

$$Y=6$$

$$6=c+m(1)+a(1)^2$$

$$6=c+m+a$$

### Second point

$$X=2$$

$$Y=11$$

$$11=c+2m+4a$$

### Third point

$$X=3$$

$$Y=18$$

$$18=c+3m+9a$$

$$6=c+m+a$$

$$11=c+2m+4a$$

$$18=c+3m+9a$$

These are three simultaneous equations.

By solving them,  
we get  
 $c=3$   $m=2$   $a=1$

So the polynomial equation becomes

$$y=3+2X+X^2$$

we use the Least Squares Method to solve three equations simultaneously.

### Applications:

Financial market trends

Scientific experiments

### Pros

- Fits curved relationships
- More flexible than linear

### Cons

- Risk of overfitting
- Harder to interpret

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

Regression analysis is a statistical method for modeling and analyzing the relationship between a dependent variable and one or more independent variables.

**Regression is used to:**

Predict future values. Understand the relationship between variables.

### Terminology:

**Dependent Variable (Y):** Variable to be predicted

**Independent Variable (X):** Variable used for prediction

**Regression Line:** Best-fit line showing the relationship

**Coefficient:** Indicates how much Y changes when X changes

# Types of regression:

1. Simple Linear Regression
2. Multiple Linear Regression
3. Polynomial Regression

Linear Regression is a supervised machine learning algorithm used to predict a continuous numerical value by establishing a linear relationship between the input features (independent variables) and the output (dependent variable).

**by fitting the best possible straight line to the observed data.**

linear because the model assumes that the relationship between the input variable(s) and the output variable can be **represented by a straight line.**

### Examples:

Student marks

House prices

Salary

Temperature

data about students' study hours and marks. Can we predict the marks of a new student?"

| Student | Hours Studied | Marks |
|---------|---------------|-------|
| A       | 1             | 40    |
| B       | 2             | 50    |
| C       | 3             | 60    |
| D       | 4             | 70    |
| E       | 5             | 80    |

If a new student studies for 6 hours, how many marks do you think they will get?

**observed a pattern.**

1 hour → 40 marks


2 hours → 50 marks

.

.

Every extra hour increases the marks by 10.

Machine learns this pattern from the training data and predicts the marks for a new student.



Linear Regression equation:

$$Y=mX+c$$

**X = Hours studied (input).** This is the input given to the model.

**Y = Predicted marks (output) :** This is the value we want to predict of a new student.

**M (slope)** = How much Y changes when X increases by one unit.(Increase in marks for every additional hour )

**C (Intercept)** = the starting value, i.e., the predicted marks when X=0.

Every additional hour increases the marks by 10.

$$m=10$$

Find the Starting Value (c)

$$Y=mX+c$$

first student studied 1 hour and scored 40 marks.

$$40 = 10(1) + c$$

$$c = 30$$

$$Y = 10X + 30$$
 final regression equation for the given dataset

What is the initial value of Y when student studies 0 hours Substitute X  
= 0 (initial hour)

$$Y = 10(0) + 30$$

$$Y = 30$$
 (not required this)

How can we predict the marks for new students, if he study 6 hours?

We use this equation of the line.

$$Y=10X+30$$

Substitute

$$X = 6$$

$$Y=10(6)+30$$

$$Y=90$$

Marks

90 |

80 |

70 |

60 |

50 |

40 |

+

1

2

3


4

5

6

?

Hours



The straight line acts like a rule.

$$Y=10X+30$$

Now it can predict any value.

ARTIFICIAL INTELLIGENCE


Linear Regression



| Hours | Marks |
|-------|-------|
| 1     | 42    |
| 2     | 47    |
| 3     | 61    |
| 4     | 69    |
| 5     | 78    |

## What if the Data Is Not Perfect?

Real-world data is not perfect.  
The points are not perfectly aligned



In real-world data, all points do not lie exactly on a straight line. Therefore, Linear Regression finds the line that is closest to all the data points.

This is called the best-fit line.

How does the machine find the best-fit line?

It uses the **Least Squares Method**.

**Least Squares formulas mathematically** calculate the best values of

$m$

$c$

instead of guessing them.

This method calculates the prediction error for the line and selects the line with the minimum Sum of Squared Errors (SSE)."

Using these two formulas.

$$m = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$$

$$c = \bar{y} - m\bar{x}$$

## Step 2: Calculate the Mean

Mean of X

$$\bar{x} = \frac{1 + 2 + 3 + 4 + 5}{5} = \frac{15}{5} = 3$$

Mean of Y

$$\bar{y} = \frac{42 + 47 + 61 + 69 + 78}{5} = \frac{297}{5} = 59.4$$

## Step 3: Prepare the Table

| $x_i$ | $y_i$ | $x_i - \bar{x}$ | $y_i - \bar{y}$ | $(x_i - \bar{x})(y_i - \bar{y})$ | $(x_i - \bar{x})^2$ |
|-------|-------|-----------------|-----------------|----------------------------------|---------------------|
| 1     | 42    | -2              | -17.4           | 34.8                             | 4                   |
| 2     | 47    | -1              | -12.4           | 12.4                             | 1                   |
| 3     | 61    | 0               | 1.6             | 0                                | 0                   |
| 4     | 69    | 1               | 9.6             | 9.6                              | 1                   |
| 5     | 78    | 2               | 18.6            | 37.2                             | 4                   |

Now calculate the totals.

$$\sum (x_i - \bar{x})(y_i - \bar{y}) = 34.8 + 12.4 + 0 + 9.6 + 37.2 = 94$$

$$\sum (x_i - \bar{x})^2 = 4 + 1 + 0 + 1 + 4 = 10$$

### Step 4: Calculate the Slope (m)

Formula:

$$m = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2}$$

Substitute the values:

$$m = \frac{94}{10} = 9.4$$

So,

$$m = 9.4$$

## Step 5: Calculate the Intercept (c)

Formula:

$$c = \bar{y} - m\bar{x}$$

Substitute the values:

$$c = 59.4 - (9.4 \times 3)$$

$$c = 59.4 - 28.2$$

$$c = 31.2$$

So,

$$c = 31.2$$

## Step 6: Final Linear Regression Equation

Substitute  $m$  and  $c$  into

$$Y = mX + c$$

$$Y = 9.4X + 31.2$$

## Step 7: Verify the Equation

For 3 hours:

$$Y = 9.4(3) + 31.2$$

$$Y = 28.2 + 31.2$$

$$Y = 59.4$$

Actual marks = **61**

Predicted marks = **59.4**

The prediction is close to the actual value.

This equation is then used to predict the marks of new students based on their study hours."

$$Y=9.4X+31.2$$

For 6 hours:

$$Y = 9.4(6) + 31.2$$

$$Y = 56.4 + 31.2$$

$$Y = 87.6$$

Predicted marks  $\approx$  88

### Pros

- Simple & easy to understand
- Fast to train and implement
- Works well for linear relationships
- Useful for prediction and forecasting
- Easy interpretability of coefficients

### Cons

- Computationally expensive for large datasets
- Poor performance when data is noisy/overlapping
- Less interpretable

# Dataset 1: Salary Prediction (Recommended)

Problem: Predict an employee's salary based on years of experience.

| Years of Experience (X) | Salary (₹ in Lakhs) (Y) |
|-------------------------|-------------------------|
| 1                       | 3                       |
| 2                       | 5                       |
| 3                       | 7                       |
| 4                       | 9                       |
| 5                       | 11                      |

### Dataset 2: House Price Prediction

| House Area (sq.ft.) | Price (₹ Lakhs) |
|---------------------|-----------------|
| 600                 | 28              |
| 800                 | 36              |
| 1000                | 47              |
| 1200                | 54              |

## Evaluation Metrics in Linear Regression

After training a Linear Regression model, we obtain a regression equation such as

$$Y = mX + c$$

This equation is used to predict the output for new data.

the predicted values may not exactly match the actual values because real-world data contains noise and variations.

it is necessary to evaluate how well the regression model performs.

## Evaluation Metrics in Linear Regression

To evaluate the model, we compare the actual values with the predicted values and calculate various evaluation metrics.

The commonly used evaluation metrics are:

1. Residual (Error)
2. Mean Squared Error (MSE)
3. Root Mean Squared Error (RMSE)
4. R-Square ( $R^2$  Score)

These metrics help determine whether the regression model is accurate and reliable.

## Residual

- Actual – Predicted
- Also represented as y-yhat

$$\text{Residual} = \text{Actual Value} - \text{Predicted Value}$$

$$\text{Residual} = Y - \hat{Y}$$

Residual indicates how far the model's prediction is from the actual value.

1. A small residual means the prediction is close to the actual value.
2. A large residual means the prediction is inaccurate.

Therefore, a good regression model should have small residual values

## Mean Squared Error (MSE)

Mean Squared Error (MSE) is the average of the squared residuals.

Residuals can be positive or negative.

If we simply add them, they may cancel each other. (-3 and +3) errors exist, the total becomes zero.

To avoid this problem, we square every residual before calculating the average.

### Interpretation

Smaller MSE → Better regression model

Larger MSE → Poor regression model

$$MSE = \frac{1}{n} \sum (Y - \hat{Y})^2$$

### Root Mean Squared Error (RMSE)

Root Mean Squared Error (RMSE) is the square root of the Mean Squared Error.

### Interpretation

Lower RMSE → Better prediction accuracy

Higher RMSE → Poor prediction accuracy

$$RMSE = \sqrt{MSE}$$

## R-Square ( $R^2$ Score)

**R-Square (Coefficient of Determination) measures how well the regression model explains the variation in the dependent variable.**

Unlike MSE and RMSE, which measure prediction error,  $R^2$  measures how well the model fits the data.

A value closer to 1 indicates that the model explains most of the variation in the data.

$$R^2 = 1 - \frac{SS_{Residual}}{SS_{Total}}$$

$$R^2 = 1 - \frac{\sum(Y - \hat{Y})^2}{\sum(Y - \bar{Y})^2}$$

where

$SS_{Residual}$  = Residual Sum of Squares (Unexplained Variation)

$SS_{Total}$  = Total Sum of Squares (Total Variation)

# ARTIFICIAL INTELLIGENCE

## Regression Model Evaluation

| GIVEN DATA (Training / Test Data) |                  |
|-----------------------------------|------------------|
| X (Hours)                         | Y (Actual Marks) |
| 1                                 | 42               |
| 2                                 | 47               |
| 3                                 | 61               |
| 4                                 | 69               |
| 5                                 | 78               |

### REGRESSION EQUATION (Best Fit Line)

$$\hat{Y} = 9.4X + 31.2$$

### MEAN OF ACTUAL VALUES

$$\bar{Y} = \frac{42+47+61+69+78}{5} = \frac{297}{5} = 59.4$$

| X<br>(Hours)       | Y<br>(Actual<br>Marks) | $\hat{Y}$<br>(Predicted Marks)<br>$= 9.4X + 31.2$ | Residual<br>( $Y - \hat{Y}$ )<br>(e) | Squared<br>Residual<br>( $Y - \hat{Y}$ ) <sup>2</sup> | $Y - \bar{Y}$<br>(Deviation from Mean) | ( $Y - \bar{Y}$ ) <sup>2</sup> |
|--------------------|------------------------|---------------------------------------------------|--------------------------------------|-------------------------------------------------------|----------------------------------------|--------------------------------|
| 1                  | 42                     | 40.6                                              | 1.4                                  | 1.96                                                  | -17.4                                  | 302.76                         |
| 2                  | 47                     | 50.0                                              | -3.0                                 | 9.00                                                  | -12.4                                  | 153.76                         |
| 3                  | 61                     | 59.4                                              | 1.6                                  | 2.56                                                  | 1.6                                    | 2.56                           |
| 4                  | 69                     | 68.8                                              | 0.2                                  | 0.04                                                  | 9.6                                    | 92.16                          |
| 5                  | 78                     | 78.2                                              | -0.2                                 | 0.04                                                  | 18.6                                   | 345.96                         |
| TOTAL ( $\Sigma$ ) | 297                    | 296.0                                             | 0.0                                  | <b>13.60</b>                                          | 0.0                                    | <b>897.20</b>                  |

$$MSE = \frac{\sum (Y - \hat{Y})^2}{n} = \frac{13.60}{5} = 2.72$$

$$RMSE = \sqrt{MSE} = \sqrt{2.72} = 1.65$$

$$\begin{aligned}
 R^2 &= 1 - \frac{SS_{Residual}}{SS_{Total}} \\
 &= 1 - \frac{13.60}{897.20} = 1 - 0.01516 = \mathbf{0.985}
 \end{aligned}$$

| Metric | Value | Interpretation        |
|--------|-------|-----------------------|
| MSE    | 2.72  | Lower is better       |
| RMSE   | 1.65  | Lower is better       |
| $R^2$  | 0.985 | Closer to 1 is better |

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

**Support Vector Machine (SVM)** is a supervised machine learning algorithm used for classification and regression problems.

It works by finding the optimal hyperplane that separates different classes with the maximum possible margin.


The main **objective of SVM** is to maximize the distance (margin) between the separating hyperplane and the nearest data points from each class, resulting in better generalization and classification accuracy.

### Tom M. Mitchell

"A Support Vector Machine is a learning algorithm that constructs a hyperplane or set of hyperplanes in a high-dimensional space for classification or regression."

- Imagine you have two groups of students:
  - Some prefer Cricket
  - Some prefer Football
- You want to separate them in a playground.
- You draw a line such that:
  - Cricket lovers are on one side
  - Football lovers are on the other side
- This is what SVM does – it finds the best dividing line (hyperplane)

- Works by **finding a hyperplane that best separates classes**
- Goal: **Maximize the margin** between the classes and the dividing line
- Support vectors = Data points closest to the line



## Margin

Margin means empty space around the hyperplane. The blank space above and below the line is called the Margin.

**SVM always tries to make this margin as wide as possible.**

## A larger margin means:

Better separation

Less confusion

Better prediction on new data

Less overfitting

That is why SVM looks for the maximum margin hyperplane.

### Small Road



Vehicles may collide.

### Wide Road



Safer.


- **Hyperplane** → Best separating boundary
- **Margin** → Distance between hyperplane & closest data points
- **Support Vectors** → Points that lie on margin boundaries and influence the hyperplane

## Working of SVM


1. Plot data points of two classes
2. Find a line (hyperplane) that separates the classes
3. Choose the one with the maximum margin (largest distance between nearest points of classes)
4. Use this hyperplane to classify new data



=====



the points nearest to the hyperplane called Support Vectors. The other points far away don't affect the hyperplane



### The SVM algorithm:

1. Looks at all training points.
2. Finds many possible separating lines.
3. Calculates the margin for each line.
4. Chooses the line with the **maximum margin**.
5. Identifies the **support vectors** (closest points).
6. Stores the decision boundary as the trained model.

**SVM find the (shortest) distance** from each data point to the hyperplane.

What is the hyperplane?

**With one feature:**

Hyperplane is point/threshold

**With 2 features:**

The hyperplane is a **line**. Represent in 2D

**With three feature:**

the data exists in **3-dimensional space**. SVM decision boundary becomes a **plane**.

**More than three feature means Hyperplane.**

**Support Vectors**

The distance between the boundary and the closest points is the **margin**.

Applying SVM to the following dataset


| Player | Avg Runs | Avg. Wickets | Class   |
|--------|----------|--------------|---------|
| P1     | 50       | 1            | Batsman |
| P2     | 60       | 0            | Batsman |
| P3     | 45       | 2            | Batsman |
| P4     | 10       | 3            | Bowler  |
| P5     | 15       | 4            | Bowler  |
| P6     | 20       | 5            | Bowler  |

1. How machine will train. The machine starts looking for relationships.
2. Avg runs and Avg Wickets are the two features.


1. Low Runs, High Wickets Mostly Bowlers
2. High Runs, Low Wickets Mostly Batsmen

1. For SVM, the machine finds the best hyperplane that separates the two classes.
2. After training, the machine stores the decision boundary.
3. Suppose a new player comes.
4. The model checks: "Which side of the hyperplane is this player?"

## Cricket Dataset: Batsman vs Bowler



## SVM: Batsman vs Bowler (Runs vs Wickets)



The two points **closest to the other class** are:

- Bowler → **P6 = (20, 5)**
- Batsman → **P3 = (45, 2)**

These are the **support vectors**

using a **simplified SVM** where **Runs alone determines the separation**, the decision boundary is exactly halfway between their Run values:

$$\begin{aligned}\text{Decision Boundary} &= \frac{\text{Bowler support vector} + \text{Batsman support vector}}{2} \\ &= \frac{20 + 45}{2} \\ &= \boxed{32.5}\end{aligned}$$


So the decision boundary is:

$$\boxed{x = 32.5}$$

The boundary is equally far from both support vectors:

$$32.5 - 20 = 12.5$$


$$45 - 32.5 = 12.5$$



Here consider the only one feature : runs.

- Line roughly between the two groups:
  - Above line → Bowler
  - Below line → Batsman
- Identify the support vectors → Closest points from each class that “support” the margin:
  - Batsman near the line (eg, P3)
  - Bowler near the line (eg, P4)


- Suppose a new player has: **Runs = 30, Wickets = 3**
- Check which side of the line the player lies:
  - If above → Bowler
  - If below → Batsman”



# What if there are 3 classes?

Suppose we have:

- Apple
- Mango
- Orange




**Can one hyperplane separate all three classes?**

SVM is a binary classifier.

That means it can classify only 2 classes at a time.

SVM creates multiple hyperplanes.

Y-axis



- Apple
- Mango
- Orange

Take only two classes at a time and ignore all the others.

One-vs-One

**Classifier 1**

Apple vs Mango

**Classifier 2**

Apple vs Orange

**Classifier 3**

Mango vs Orange

Suppose we build **Apple vs Rest.** (One VS rest)

## Pros

- Works well with high-dimensional data
- Effective when there is a clear margin of separation
- Good performance for small/medium datasets
- The maximum margin principle helps SVM generalize well to unseen data, reducing overfitting.
- SVM often provides excellent performance even when the training dataset is not very large.

### Cons

- Computationally expensive for large datasets
- Poor performance when data is noisy/overlapping
- Less interpretable:
- Unlike Decision Trees, the decision-making process of SVM is not easy to understand or visualize.
- Training time increases as the dataset size and feature count grow.

# Applications

**Write the real time application where SVM algorithm is suitable**

| Student | Study Hours<br>( $X_1$ ) | Attendance<br>( $X_2$ ) | Class |
|---------|--------------------------|-------------------------|-------|
| S1      | 1                        | 50                      | Fail  |
| S2      | 2                        | 55                      | Fail  |
| S3      | 2                        | 60                      | Fail  |
| S4      | 3                        | 65                      | Pass  |
| S5      | 4                        | 70                      | Pass  |
| S6      | 5                        | 75                      | Pass  |

## Classify a student

Suppose a new student has:

- Study Hours = 4
- Attendance = 70%

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

## K-Nearest Neighbors (KNN) Algorithm

### Definition

K-Nearest Neighbors (KNN) is a supervised machine learning algorithm used for classification and regression problems. It classifies a new data instance by finding the K closest training examples in the feature space and assigning the class that appears most frequently among them.


KNN is called a **Lazy Learning Algorithm** because it does not build a model during the training phase. Instead, it stores all the training data and performs computation only when a new data point needs to be classified.

## Principle of KNN

- When a new data point is received,
- KNN: Measures the distance between the new point and every training instance.
- Finds the K nearest neighbors.
- Predicts the class based on majority voting.

K is the number of nearest neighbors that the algorithm considers before making a prediction.

## Working of KNN



## Importance of Selecting the Value of K in KNN

The value of K in the K-Nearest Neighbors (KNN) algorithm represents the number of nearest neighbors considered when predicting the class of a new data point.

Selecting an appropriate value of K is important because it directly affects the accuracy, stability, and performance of the KNN model.

### Small K (e.g., $K = 1$ )

When  $K = 1$ , the algorithm considers only the nearest neighbor. May produce incorrect predictions if the nearest neighbor is mislabeled.

### Moderate K (e.g., $K = 3$ or $K = 5$ )

The algorithm considers multiple neighbors and predicts using majority voting.

Produces more reliable predictions. Usually gives better generalization.

## Algorithm

Step 1: Store all training data.

Step 2: Choose the value of K.

Step 3: Calculate the distance between the new data point and every training instance.

The most commonly used distance is Euclidean Distance.

$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

where

$x_i$  = feature value of test sample

$y_i$  = feature value of training sample

$n$  = number of features

Step 4: Sort all distances in ascending order.

Step 5: Select the first K nearest neighbors.

Step 6: Count the number of occurrences of each class among the K neighbors.

Step 7: Assign the class with the highest frequency to the new data point.

Step 8: Display the predicted class.

**1. TRAINING DATASET**

| Patient | Glucose<br>(mg/dL) | BMI<br>(kg/m <sup>2</sup> ) | Class       |
|---------|--------------------|-----------------------------|-------------|
| P1      | 80                 | 22                          | No Diabetes |
| P2      | 85                 | 24                          | No Diabetes |
| P3      | 90                 | 26                          | No Diabetes |
| P4      | 140                | 32                          | Diabetes    |
| P5      | 150                | 34                          | Diabetes    |
| P6      | 160                | 36                          | Diabetes    |

 **No Diabetes** **Diabetes**

### NEW PATIENT (TO BE CLASSIFIED)

Glucose = 145 mg/dL


BMI = 33 kg/m<sup>2</sup>

✕ (145, 33)

Choose **K = 3**

The algorithm considers the three nearest neighbors.

## 2. DATA POINTS IN 2D SPACE



### 3. CALCULATE EUCLIDEAN DISTANCE

Distance formula:

$$d = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}$$

where

$$x_1 = 145, \quad x_2 = 33$$

| Patient | Calculation                      | Distance      |
|---------|----------------------------------|---------------|
| P1      | $\sqrt{(145-80)^2 + (33-22)^2}$  | = 65.92       |
| P2      | $\sqrt{(145-85)^2 + (33-24)^2}$  | = 60.67       |
| P3      | $\sqrt{(145-90)^2 + (33-26)^2}$  | = 55.44       |
| P4      | $\sqrt{(145-140)^2 + (33-32)^2}$ | = <b>5.10</b> |
| P5      | $\sqrt{(145-150)^2 + (33-34)^2}$ | = <b>5.10</b> |
| P6      | $\sqrt{(145-160)^2 + (33-36)^2}$ | = 15.30       |

Here:

- $x_1 = 145 \rightarrow$  Feature 1 of the test sample
- $x_2 = 33 \rightarrow$  Feature 2 of the test sample

Test patient:

$$(x_1, x_2) = (145, 33)$$

Training patient P1:

$$(y_1, y_2) = (80, 22)$$

Euclidean Distance

$$d = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2}$$

Substitute the values:

$$d = \sqrt{(145 - 80)^2 + (33 - 22)^2}$$

$$d = \sqrt{65^2 + 11^2}$$

$$d = \sqrt{4225 + 121}$$

$$d = \sqrt{4346}$$

$$\boxed{d = 65.92}$$

#### 4. SORT DISTANCES (ASCENDING ORDER)

| Rank | Patient | Distance | Class       |
|------|---------|----------|-------------|
| 1    | P4      | 5.10     | Diabetes    |
| 2    | P5      | 5.10     | Diabetes    |
| 3    | P6      | 15.30    | Diabetes    |
| 4    | P3      | 55.44    | No Diabetes |
| 5    | P2      | 60.67    | No Diabetes |
| 6    | P1      | 65.92    | No Diabetes |

### 5. SELECT $K = 3$ NEAREST NEIGHBORS

| Selected Neighbors ( $K=3$ ) | Class    |
|------------------------------|----------|
| P4 (Distance = 5.10)         | Diabetes |
| P5 (Distance = 5.10)         | Diabetes |
| P6 (Distance = 15.30)        | Diabetes |

### 6. MAJORITY VOTING

Diabetes = 3 votes      No Diabetes = 0 votes

### 7. PREDICTION

The new patient is predicted to have **DIABETES**

### Pros

- Very simple and easy to understand and implement.
- Works very well for datasets with less number of features
- KNN can be used to solve both classification (e.g., spam detection) and regression (predicting continuous values) problems.
- KNN performs well when the dataset is small and the classes are well separated.
- Few parameters:  $k$ , Distance Metric

### Cons

- To classify a new sample, KNN calculates the distance to all training samples. This becomes slow for large datasets.
- KNN stores the entire training dataset, which requires more memory.
- Incorrect or noisy training samples can influence the prediction, especially when K is small.
- As the number of features increases, distances between data points become less meaningful. This is known as the curse of dimensionality.

## Applications :

- KNN is widely used in applications where similarity between data points is important. Since KNN predicts a new instance based on its nearest neighbors, it is suitable for recommendation, classification, and pattern recognition tasks.
- 1. **Image Classification:** A new image of an animal is compared with stored images. If most similar images are cats, the new image is classified as a cat.
  - Animal identification
  - Plant species classification
  - Object recognition
- **2. Face Recognition**
  - Smartphone face unlock
  - Attendance systems
  - Security and surveillance

## Applications :

- **3. Email Spam Detection**

- Classify emails as Spam or Not Spam. A new email is compared with previously labeled emails. If most similar emails are spam, it is marked as spam.

- **4. Recommendation Systems**

- Recommend products, movies, or songs based on user preferences.
- Movie recommendations
- Music recommendations
- Online shopping

- **5. Customer Segmentation:** Group customers with similar purchasing behavior.

- **6. Document and Text Classification :**

news article is classified as:

Sports, Politics, Technology, Entertainment

by comparing it with previously labeled articles.

Practice

| Name   | Acid Durability | Strength | class |
|--------|-----------------|----------|-------|
| Type-1 | 7               | 7        | Bad   |
| Type-2 | 7               | 4        | Bad   |
| Type-3 | 3               | 4        | Good  |
| Type-4 | 1               | 4        | Good  |
| Type-5 | 3               | 7        | ??    |

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

**A Decision Tree** is a supervised machine learning algorithm that makes decisions by asking a sequence of questions (yes/No) about the input features.

makes predictions by splitting the data into smaller groups based on a sequence of conditions.

**It resembles a tree structure**, where each decision is based on a feature, and the final outcome is reached after evaluating these conditions.

## Characteristics of Decision Tree

It is a Supervised Learning algorithm.

Can be used for both Classification and Regression.

Represents decisions in a tree-like hierarchical structure.

Easy to understand and interpret.

Works with both **numerical and categorical** data.

Requires minimal data **preprocessing**.

## Simple Analogy

### 1. Doctor's Diagnosis

A doctor does not prescribe medicine after asking only one question. The doctor asks about symptoms, medical history, and test results before reaching a diagnosis.

Likewise, a Decision Tree examines different features before making a decision.

### 2. Shopping for a Laptop ?

A Decision Tree consists of three main components:

## 1. Root Node (Tree node)

The topmost node of the tree.

Represents the best feature selected for the first split.

There is only one root node in a decision tree.

## 2. Decision Node (Internal Node)

Represents a test or condition on a feature.

Splits the data into smaller subsets based on the feature values.

A tree can have multiple decision nodes.

Example: Diameter  $>$  8 cm?

---

A Decision Tree consists of three main components:

### 3. Leaf Node (Terminal Node)

Represents the final prediction or decision.


Does not split further.

Contains the class label or predicted value.

**Example:** Apple or Mango

| Training Dataset |            |        |               |                    |
|------------------|------------|--------|---------------|--------------------|
| ID               | Weight (g) | Color  | Diameter (cm) | Fruit              |
| 1                | 160        | Red    | 8             | Apple              |
| 2                | 165        | Red    | 8             | Apple              |
| 3                | 170        | Red    | 9             | Strawberry         |
| 4                | 175        | Red    | 10            | Strawberry         |
| 5                | 180        | Yellow | 9             | Mango              |
| 6                | 185        | Yellow | 10            | Mango              |
| 7                | 200        | Yellow | 11            | Other Yellow Fruit |
| 8                | 220        | Yellow | 12            | Other Yellow Fruit |

## Example 1



The Decision Tree algorithm learns from the training data and identifies patterns between the **features (Weight, Color, Diameter)** and the target label (Fruit Name).

Based on these patterns, it automatically creates a set of decision rules, which together form the Decision Tree Model.


**Model:** A trained machine learning model that can classify new fruits

# Decision Tree Construction Algorithm

1. Start with the complete fruit dataset.
2. Select the most informative feature as the Root Node.
3. Create a decision question using the selected feature.
4. Split the dataset into smaller subsets.
5. Check whether each subset contains only one fruit class.
6. If Yes, create a Leaf Node.
7. Otherwise, repeat the process using the remaining features.
8. Continue until all leaf nodes are pure or no further splitting is possible.

## Example 2


## Decision Tree for Play Tennis



| Training Dataset |            |        |               |                    |
|------------------|------------|--------|---------------|--------------------|
| ID               | Weight (g) | Color  | Diameter (cm) | Fruit              |
| 1                | 160        | Red    | 8             | Apple              |
| 2                | 165        | Red    | 8             | Apple              |
| 3                | 170        | Red    | 9             | Strawberry         |
| 4                | 175        | Red    | 10            | Strawberry         |
| 5                | 180        | Yellow | 9             | Mango              |
| 6                | 185        | Yellow | 10            | Mango              |
| 7                | 200        | Yellow | 11            | Other Yellow Fruit |
| 8                | 220        | Yellow | 12            | Other Yellow Fruit |


### STEP 1

Start with the complete dataset.  
Select the best feature as the Root Node.  
(Best feature = Color)




### STEP 2

Split the dataset based on Color.  
Check each subset.




### Legend

-  Decision Node (Question)
-  Data Subset
-  Leaf Node (Class Label)


STEP 3

Red subset is not pure (contains Apple and Strawberry). Select next best feature = Diameter and split.




STEP 4

Yellow subset is not pure (contains Mango and Other Yellow Fruit). Select next best feature = Weight and split.




**STEP 5**

All the subsets are now pure.  
Create leaf nodes with their class labels.

**STEP 6**


The final Decision Tree is ready and can be used to classify new fruits.



### Test Sample

| Feature  | Value |
|----------|-------|
| Weight   | 165 g |
| Color    | Red   |
| Diameter | 8 cm  |


### Learned Decision Tree (from Training Data)



### How Decision Tree Classifies the Test Sample

#### Step 1


Start at the root node.  
Check Color?



Test sample Color = **Red**  
So, follow the **Red** branch.

#### Step 2


Now at node: Diameter  $\leq 8$  cm?  
Check Diameter.



Test sample Diameter = 8 cm  
 $8 \leq 8 \rightarrow$  **Yes**  
So, follow the Yes branch.

#### Step 3

Follow the branch to the leaf node.



We reach the leaf node.

Thus, for the given test sample:

- Weight = 165 g
- Color = Red
- Diameter = 8 cm

The Decision Tree predicts  
the fruit as  $\rightarrow$  **Apple**

## Model Evaluation

In Machine Learning, we first divide the dataset into Training Data and Testing Data. The model learns patterns from the training data. Then we evaluate its performance using the testing data."

Our fruit dataset is very small and simple. The patterns are easy to learn. For example, almost all Red fruits are Apple and Yellow fruits are Mango. Therefore, the Decision Tree correctly classified every test sample, resulting The Decision Tree achieved 100% accuracy

## Model Evaluation

Real-world datasets are much more complex. Different classes often have overlapping features, missing values, or noisy data. Therefore, models usually make some incorrect predictions, and the accuracy is often less than 100%."

Here, Color alone cannot distinguish the fruits, so the model has to consider multiple features and may still make mistakes.

### Training Data

|   | Weight | Color  | Diameter | Fruit |
|---|--------|--------|----------|-------|
| 0 | 160    | Red    | 8        | Apple |
| 1 | 170    | Red    | 8        | Apple |
| 2 | 180    | Yellow | 9        | Mango |
| 3 | 220    | Yellow | 10       | Mango |
| 4 | 165    | Red    | 8        | Apple |
| 5 | 200    | Yellow | 9        | Mango |
| 6 | 175    | Yellow | 8        | Apple |
| 7 | 210    | Yellow | 10       | Mango |

### Testing Data

|   | Weight | Color  | Diameter | Fruit |
|---|--------|--------|----------|-------|
| 0 | 185    | Red    | 9        | Mango |
| 1 | 168    | Red    | 8        | Apple |
| 2 | 205    | Yellow | 9        | Mango |
| 3 | 190    | Yellow | 8        | Apple |

# Why do we need Model Evaluation?

Model Evaluation helps us measure how well a machine learning model performs on unseen data. It tells us whether the model's predictions are reliable enough to use in real-world applications.

It helps us answer questions like:

Is the model making correct predictions?

How many mistakes is it making?

Can we trust this model?

Is this model better than another model?

"Without model evaluation, we would not know which algorithm performs better."

There is no single machine learning algorithm that is best for every problem. Different datasets have different characteristics, so different algorithms perform differently.

## Model Evaluation?

Model Evaluation is the process of measuring how well a trained machine learning model performs on unseen (test) data.

It helps us determine whether the model makes correct predictions and whether it can be trusted for real-world applications.

## Why is Model Evaluation Important?

Measures the model's performance.

Checks how accurately the model predicts new data.

Compares different machine learning algorithms.

Helps select the best model.

Detects overfitting and underfitting.

### Advantages

- Simple and Results can be displayed as a tree diagram, making interpretation straightforward.
- Works with both numbers and categories
- Once trained, it quickly predicts the class of new data.

### Limitations

- Sensitive to small changes in the dataset..
- it performs very well on the training data but poorly on new, unseen data (Overfitting).
- Performance can be affected by imbalanced datasets
- Large trees become difficult to interpret.

## New Fruit

| Weight | Color | Fruit |
|--------|-------|-------|
| 185    | Red   | Mango |

The model predicts:

Red → Apple

But the correct answer is:

Mango

The model **memorized** the training pattern and could not **generalize** to a new case.

This is **Overfitting**.

**Underfitting** means the model has not learned enough from the training data.

**Overfitting** means the model has learned too much from the training data, including mistakes or noise.

# Evaluation Metric:

Accuracy

Precision

Recall

F1-Score

# Confusion Matrix

A Confusion Matrix is a table used to evaluate the performance of a classification model.

It compares the model's predicted labels with the actual labels and shows how many predictions are correct and how many are incorrect.

It tells us where the model is making correct predictions and where it is making mistakes.

Decision Tree classifies fruits. After testing, we obtain:

| Test Dataset |              |                 |
|--------------|--------------|-----------------|
| Fruit ID     | Actual Fruit | Predicted Fruit |
| F1           | Apple        | Apple           |
| F2           | Mango        | Mango           |
| F3           | Apple        | Mango           |
| F4           | Mango        | Mango           |
| F5           | Apple        | Apple           |

Can you answer these questions?

- Which fruit was classified incorrectly?
- How many Apples were predicted correctly?
- How many Mangoes were predicted correctly?
- What type of mistakes did the model make?

## Confusion Matrix for the Fruit Example

| Actual \ Predicted | Apple | Mango |
|--------------------|-------|-------|
| Apple              | 2     | 1     |
| Mango              | 0     | 2     |

## Interpretation

2 Apples were correctly classified as Apple.

2 Mangoes were correctly classified as Mango.

1 Apple was incorrectly classified as Mango.

No Mango was incorrectly classified as Apple.

Thus, the model made 4 correct predictions and 1 incorrect prediction.

## Confusion Matrix for the Fruit Example

| Actual \ Predicted | Apple | Mango |
|--------------------|-------|-------|
| Apple              | 2     | 1     |
| Mango              | 0     | 2     |

## General Confusion Matrix (Binary Classification)

|                 | Predicted Positive  | Predicted Negative  |
|-----------------|---------------------|---------------------|
| Actual Positive | True Positive (TP)  | False Negative (FN) |
| Actual Negative | False Positive (FP) | True Negative (TN)  |

where one class is designated as the **positive class** and the other as the **negative class**. binary classification

The matrix shows where the model is "confused."

## Confusion Matrix

|              |       | Predicted Fruit |           |          |
|--------------|-------|-----------------|-----------|----------|
|              |       | Apple           | Mango     | Total    |
| Actual Fruit | Apple | 2<br>(TP)       | 1<br>(FN) | 3        |
|              | Mango | 0<br>(FP)       | 2<br>(TN) | 2        |
|              | Total | 2               | 3         | <u>5</u> |

**True Positive (TP):** The model correctly predicts the positive class.

**True Negative (TN):** The model correctly predicts the negative class.

**False Positive (FP):** The model predicts Positive, but the actual class is Negative.

**False Negative (FN):** The model predicts Negative, but the actual class is Positive.

**1. Accuracy:** Accuracy is the percentage of total predictions that are correctly classified by the model.

Or

Accuracy is the proportion of correctly classified instances among all instances.

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

**2. Precision:** Precision is the proportion of correctly predicted positive samples out of all samples predicted as positive.

Or

Precision measures how many of the instances predicted as positive are actually positive.

$$\text{Precision} = \frac{TP}{TP + FP}$$

**3. Recall (Sensitivity or True Positive Rate) :** Recall is the proportion of actual positive samples that are correctly identified by the model.

Or

Recall measures how many of the actual positive instances are correctly identified by the classifier.

$$\text{Recall} = \frac{TP}{TP + FN}$$

**F1-Score:** F1-Score is the harmonic mean of Precision and Recall. It provides a balanced measure when both Precision and Recall are important.

$$\text{F1-Score} = \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

or equivalently,

$$\text{F1-Score} = \frac{2TP}{2TP + FP + FN}$$

Compare the prediction with the actual value.

## Confusion Matrix

|              |       | Predicted Fruit |           |          |
|--------------|-------|-----------------|-----------|----------|
|              |       | Apple           | Mango     | Total    |
| Actual Fruit | Apple | 2<br>(TP)       | 1<br>(FN) | 3        |
|              | Mango | 0<br>(FP)       | 2<br>(TN) | 2        |
|              | Total | 2               | 3         | <u>5</u> |

TP (True Positive) = 2

TN (True Negative) = 2

FP (False Positive) = 0

FN (False Negative) = 1

- Accuracy =  $(TP + TN) / \text{Total}$   
=  $(2 + 2) / 5 = \mathbf{0.80} \rightarrow \mathbf{80\%}$
- Precision (Apple) =  $TP / (TP + FP)$   
=  $2 / (2 + 0) = \mathbf{1.00} \rightarrow \mathbf{100\%}$
- Recall (Apple) =  $TP / (TP + FN)$   
=  $2 / (2 + 1) = \mathbf{0.67} \rightarrow \mathbf{66.67\%}$
- F1-Score (Apple) =  $2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$   
=  $2 \times (1.00 \times 0.67) / (1.00 + 0.67) = 0.80 \rightarrow \mathbf{80\%}$

Rows represent Actual classes and Columns represent Predicted classes.  
Positive class (for this example) = Apple

## Applications:

### Medical Diagnosis

Predict diseases based on symptoms such as fever, cough, blood pressure, and test results.

### Bank Loan Approval

Decide whether to approve or reject a loan using income, credit score, employment status, and repayment history.

### Email Spam Detection

Classify emails as Spam or Not Spam using keywords, sender information, and attachments.

### Agriculture

Recommend suitable crops based on soil type, rainfall, temperature, and season.

### Quality Control

Detect defective and non-defective products in manufacturing using measurements and inspection results.

### Exercise

A machine learning model was developed to predict whether an email is Spam (Positive) or Not Spam (Negative).

When tested on 100 emails, the results were as follows:

- The model correctly predicted 40 spam emails as spam
- It correctly predicted 50 non-spam emails as not spam
- It incorrectly marked 5 non-spam emails as spam
- It failed to detect 5 spam emails and marked them as not spam
- Construct Confusion matrix and calculate accuracy, precision, recall.

## Training

The system is given thousands of emails.

### Email

"Win ₹1 Crore!"

"Meeting at 2 PM"

"Claim your prize"

"Project Report"

### Label

Spam

Not Spam

Spam

Not Spam

The machine learns patterns:

- "Prize"
- "Free"
- "Lottery"

are common in spam emails.

## Testing

New Email:

"Congratulations! Claim your free reward."

The machine compares it with what it learned during training and predicts:

➡ **Spam**

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications

ARTIFICIAL INTELLIGENCE

EL- DT : Disease Prediction (Medium)



| Age | Fever   | Cough   | Sore Throat | Contact Risk | Oxygen | COVID Positive |
|-----|---------|---------|-------------|--------------|--------|----------------|
| 25  | 1 (Yes) | 1 (Yes) | 0 (No)      | High         | 92     | 1 (Positive)   |
| 40  | 0 (No)  | 1 (Yes) | 1 (Yes)     | Low          | 97     | 0 (Negative)   |
| 35  | 1 (Yes) | 0 (No)  | 1 (Yes)     | Medium       | 96     | 0 (Negative)   |
| 60  | 1 (Yes) | 1 (Yes) | 0 (No)      | High         | 89     | 1 (Positive)   |
| 50  | 0 (No)  | 0 (No)  | 0 (No)      | Low          | 98     | 0 (Negative)   |

ARTIFICIAL INTELLIGENCE

EL- DT : Disease Prediction



← Features → Label

| Age | Fever   | Cough   | Sore Throat | Contact Risk | Oxygen | COVID Positive |
|-----|---------|---------|-------------|--------------|--------|----------------|
| 25  | 1 (Yes) | 1 (Yes) | 0 (No)      | High         | 92     | 1 (Positive)   |
| 40  | 0 (No)  | 1 (Yes) | 1 (Yes)     | Low          | 97     | 0 (Negative)   |
| 35  | 1 (Yes) | 0 (No)  | 1 (Yes)     | Medium       | 96     | 0 (Negative)   |
| 60  | 1 (Yes) | 1 (Yes) | 0 (No)      | High         | 89     | 1 (Positive)   |
| 50  | 0 (No)  | 0 (No)  | 0 (No)      | Low          | 98     | 0 (Negative)   |

Train data

Test data

Numerical Feature


Categorical Feature

Model : Learn from data

If cough =Yes & Contact Risk=high, then covid=Positive

Prediction

Age=30, Fever=yes, cough=yes, sore throat=no, contact risk=medium, oxygen=88 , then covid=?



### Test sample:

Patient: Age 45, Fever = Yes, Cough = Yes,  
Oxygen = 92%, Contact Risk = Medium


### Tree check:

If Cough=yes

->Contact Risk=high ->No

-> Oxygen <94 -> Yes ->

**Covid=Positive**



# ARTIFICIAL INTELLIGENCE

## EL- DT : Disease Prediction


| ID | Age | Fever | Cough | SoreThroat | ContactRis k | Oxygen | Covid    |
|----|-----|-------|-------|------------|--------------|--------|----------|
| T1 | 65  | Yes   | Yes   | Yes        | High         | 89     | Positive |
| T2 | 45  | Yes   | Yes   | No         | Medium       | 92     | Positive |
| T3 | 25  | No    | No    | No         | Low          | 98     | Negative |
| T4 | 52  | Yes   | No    | No         | Low          | 95     | Negative |
| T5 | 36  | Yes   | Yes   | Yes        | High         | 90     | Positive |
| T6 | 29  | No    | Yes   | No         | Medium       | 97     | Negative |
| S1 | 30  | Yes   | Yes   | No         | Medium       | 88     | Positive |
| S2 | 70  | Yes   | No    | No         | Low          | 91     | Negative |
| S3 | 40  | No    | No    | No         | High         | 96     | Negative |
| S4 | 55  | Yes   | Yes   | Yes        | High         | 94     | Positive |

## DECISION TREE (DT)

(Target Variable: Covid)

**Attributes Used:**

1. Fever
2. ContactRisk
3. Oxygen



# ARTIFICIAL INTELLIGENCE

## EL- DT : Disease Prediction

| ID | Age | Fever | Cough | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|-------|------------|-------------|--------|----------|
| T1 | 65  | Yes   | Yes   | Yes        | High        | 89     | Positive |
| T2 | 45  | Yes   | Yes   | No         | Medium      | 92     | Positive |
| T3 | 25  | No    | No    | No         | Low         | 98     | Negative |
| T4 | 52  | Yes   | No    | No         | Low         | 95     | Negative |
| T5 | 36  | Yes   | Yes   | Yes        | High        | 90     | Positive |
| T6 | 29  | No    | Yes   | No         | Medium      | 97     | Negative |

**Take Next feature**

**Predict=Negative**

Both records belong to same class label So it is Pure. Put Label as leaf

Among 4 records of Yes, 3 belong to Positive label and 1 belongs to Negative – So impure. Now Grow this side tree with the next feature. This branch will take only 4 records in the next branch

# ARTIFICIAL INTELLIGENCE

## EL- DT : Disease Prediction

| ID | Age | Fever | Cough | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|-------|------------|-------------|--------|----------|
| T1 | 65  | Yes   | Yes   | Yes        | High        | 89     | Positive |
| T2 | 45  | Yes   | Yes   | No         | Medium      | 92     | Positive |
| T3 | 25  | No    | No    | No         | Low         | 98     | Negative |
| T4 | 52  | Yes   | No    | No         | Low         | 95     | Negative |
| T5 | 36  | Yes   | Yes   | Yes        | High        | 90     | Positive |
| T6 | 29  | No    | Yes   | No         | Medium      | 97     | Negative |

=yes

=No

| ID | Age | Fever | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|------------|-------------|--------|----------|
| T1 | 65  | Yes   | Yes        | High        | 89     | Positive |
| T2 | 45  | Yes   | No         | Medium      | 92     | Positive |
| T5 | 36  | Yes   | Yes        | High        | 90     | Positive |
| T6 | 29  | No    | No         | Medium      | 97     | Negative |

| ID | Age | Fever | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|------------|-------------|--------|----------|
| T3 | 25  | No    | No         | Low         | 98     | Negative |
| T4 | 52  | Yes   | No         | Low         | 95     | Negative |

# ARTIFICIAL INTELLIGENCE

## EL- DT : Disease Prediction

=yes


| ID | Age | Fever | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|------------|-------------|--------|----------|
| T1 | 65  | Yes   | Yes        | High        | 89     | Positive |
| T2 | 45  | Yes   | No         | Medium      | 92     | Positive |
| T5 | 36  | Yes   | Yes        | High        | 90     | Positive |
| T6 | 29  | No    | No         | Medium      | 97     | Negative |

=High

| ID | Age | Fever | SoreThroat | Oxygen | Covid    |
|----|-----|-------|------------|--------|----------|
| T1 | 65  | Yes   | Yes        | 89     | Positive |
| T5 | 36  | Yes   | Yes        | 90     | Positive |

=Medium

| ID | Age | Fever | SoreThroat | Oxygen | Covid    |
|----|-----|-------|------------|--------|----------|
| T2 | 45  | Yes   | No         | 92     | Positive |
| T6 | 29  | No    | No         | 97     | Negative |



# ARTIFICIAL INTELLIGENCE

## EL- DT : Disease Prediction

| ID | Age | Fever | Cough | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|-------|------------|-------------|--------|----------|
| T1 | 65  | Yes   | Yes   | Yes        | High        | 89     | Positive |
| T2 | 45  | Yes   | Yes   | No         | Medium      | 92     | Positive |
| T3 | 25  | No    | No    | No         | Low         | 98     | Negative |
| T4 | 52  | Yes   | No    | No         | Low         | 95     | Negative |
| T5 | 36  | Yes   | Yes   | Yes        | High        | 90     | Positive |
| T6 | 29  | No    | Yes   | No         | Medium      | 97     | Negative |
| S1 | 30  | Yes   | Yes   | No         | Medium      | 88     | Positive |
| S2 | 70  | Yes   | No    | No         | Low         | 91     | Negative |
| S3 | 40  | No    | No    | No         | High        | 96     | Negative |
| S4 | 55  | Yes   | Yes   | Yes        | High        | 94     | Positive |

| ID | Age | Fever | Cough | SoreThroat | ContactRisk | Oxygen | Covid    |
|----|-----|-------|-------|------------|-------------|--------|----------|
| S1 | 30  | Yes   | Yes   | No         | Medium      | 88     | Positive |
| S2 | 70  | Yes   | No    | No         | Low         | 91     | Negative |
| S3 | 40  | No    | No    | No         | High        | 96     | Negative |
| S4 | 55  | Yes   | Yes   | Yes        | High        | 94     | Positive |

- For all these 4 records, which we kept aside, has to be used for testing. It is as good as feeding the medicine to the rats at the research lab. If the rat survives, the medicine is good. If not, go back and modify the medicine composition.

**Sklearn:** means **Scikit-learn**.

It is a popular Python library used for **Machine Learning**.

It provides algorithms and functions for:

- Decision Tree
- KNN
- SVM
- Linear Regression
- Train-test splitting
- Accuracy
- Confusion matrix
- etc.

model\_selection is a **module inside Scikit-learn**.

It contains tools used for selecting, evaluating, and preparing machine-learning models.

### **seaborn**

`seaborn` is a **Python visualization library**.

It is mainly used to create attractive statistical graphs, such as:

- Heatmaps
- Bar charts
- Histograms
- Box plots
- Scatter plots

In your program, Seaborn is being used to create the **confusion matrix heatmap**.

```
X_train, X_test, y_train, y_test = train_test_split( X, y,  
test_size=0.30, random_state=42 )
```

The dataset is divided into:

**Training data**

Used to **teach the Decision Tree.**

**Testing data**

Used to **check how well the trained tree performs on unseen data.**

```
test_size=0.30
```

This means:

- 70% → training data
- 30% → testing data

If dataset has **16 records**:

$$16 \times 0.30 = 4.8$$

Since we cannot have 4.8 rows, it becomes **5 test samples**.

```
random_state=42 mean?
```

`train_test_split()` randomly selects the records for training and testing.

If we don't specify `random_state`, the split can change every time we run the program.

### Variable

`X_train`

`X_test`

`y_train`

`y_test`

### Meaning

Features used for training

Features used for testing

Actual classes for training

Actual classes for testing

```
encoder = LabelEncoder()
```

ContactRisk contains text. Machine-learning algorithms need numerical input, so LabelEncoder converts them to numbers.

LabelEncoder determines the mapping from the labels. You can see the actual mapping using:

```
print(encoder.classes_)
```

### Select the features

```
X = data.drop("COVID", axis=1)
```

Here we create **X**, which contains the input features.

We remove the COVID column because COVID is the label we want the model to predict.

### Select the target

```
y = data["COVID"]
```

y contains the **output/target/class**.

X → Symptoms/features

y → COVID class

### Create the Decision Tree

```
model = DecisionTreeClassifier(random_state=42)
```

This creates the Decision Tree model.

### Train the Decision Tree

```
model.fit(X_train, y_train)
```

This is where the actual learning happens.

The model receives:

X\_train → Features

y\_train → Correct answers

The Decision Tree tries to find rules that separate the classes.

### Make predictions

```
y_pred = model.predict(X_test)
```

Now we give the **unseen test data** to the trained Decision Tree.

# ARTIFICIAL INTELLIGENCE

---


## ML Supervised Learning


**Dr. Rajani S**

Department of Computer Applications

Supervised Learning is a type of Machine Learning in which the model is trained using labeled data. Each training example consists of input features and a corresponding target (output) value. The objective is to learn the relationship between the inputs and outputs so that the model can accurately predict the output for new, unseen data.

**Supervised Learning is broadly divided into two main types based on the nature of the output (target variable).**






# 1. Classification

## Definition

Classification is a supervised learning task where the output is a category or class label. The model learns from labeled data and predicts which class a new input belongs to.

## Examples

- 📧 Spam Email Detection → Spam / Not Spam
- 🐱 Animal Classification → Cat, Dog, Rabbit, Bird
- 🏥 Disease Diagnosis → Diabetic / Non-Diabetic



## Classification Uses Discrete (Categorical) Output

A Discrete Variable can take specific, countable values. It cannot have decimal values between two integers.

### Examples

Number of Students = 40

Number of Cars = 25

Number of Employees = 150

Number of Books = 12

## 2. Regression

Regression is a supervised learning task where the output is a continuous numerical value. The model predicts a number rather than a category.

### Examples

House Price Prediction → ₹75,00,000

Temperature Prediction → 32°C

Stock Price Prediction → ₹1,250

Salary Prediction → ₹8,50,000/year

Rainfall Prediction → 120 mm

**Regression** uses A **Continuous Variable** can take any value within a range, including decimal values.

### Examples

House Price

Height


Weight

Temperature

Salary

Distance

Rainfall



the goal is to predict a continuous value (house price) based on the size of the house (square feet).

The house size (independent variable) is used to predict the price (Dependent variable).

Regression analysis is a statistical method for modeling and analyzing the relationship between a dependent variable and one or more independent variables.

# Difference between Classification and Regression

| Feature              | Classification                                     | Regression                                        |
|----------------------|----------------------------------------------------|---------------------------------------------------|
| Definition           | Predicts a <b>category</b> or class label.         | Predicts a <b>continuous</b> numerical value.     |
| Output Type          | Discrete (Categorical)                             | Continuous (Numerical)                            |
| Target Variable      | Discrete Variable                                  | Continuous Variable                               |
| Goal                 | Assign an input to a predefined class or category. | Estimate or predict a numerical value.            |
| Question Answered    | "Which class?"                                     | "How much?" or "What value?"                      |
| Input-Output Mapping | Maps input (X) to a <b>discrete</b> output (Y).    | Maps input (X) to a <b>continuous</b> output (Y). |
| Output Examples      | Spam / Not Spam, Cat / Dog, Pass / Fail            | ₹25,00,000, 35.5°C, ₹1,250                        |

# Difference between Classification and Regression

|                     |                                                                |                                                                       |
|---------------------|----------------------------------------------------------------|-----------------------------------------------------------------------|
| Real-World Examples | Email Spam Detection, Animal Classification, Disease Diagnosis | House Price Prediction, Temperature Prediction, Salary Prediction     |
| Common Algorithms   | Decision Tree, k-NN, SVM, Logistic Regression                  | Linear Regression, Decision Tree Regression, Random Forest Regression |
| Evaluation Metrics  | Accuracy, Precision, Recall, F1-Score                          | MAE, MSE, RMSE, R <sup>2</sup> Score                                  |

Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), R<sup>2</sup> Score

**"Where does the model get this labeled data?"**

### **Dataset**

- A dataset is simply a collection of data used to train and test a Machine Learning model.
- Think of it as a spreadsheet.
- Rows = Patients
- Columns = Information about each patient

FEATURES (X)

| Weight(g) | Color  | Diameter(cm) | Fruit Name (LABEL) |
|-----------|--------|--------------|--------------------|
| 160       | Red    | 8            | Apple              |
| 180       | Yellow | 9            | Mango              |
| 170       | Red    | 8            | Apple              |
| 220       | Yellow | 10           | Mango              |
| 165       | Red    | 8            | ?                  |

↑  
Testing Data

First 4 rows → Training Data

Model learns: If Color = Red and Diameter ≈ 8 cm

↓  
Fruit = Apple

Prediction

Weight =165  
Color = Red  
Diameter =8

↓

Fruit = Apple

### Features (Input Variables)

- The attributes/independent variables used to make predictions
- The variables can be:
  - Numerical: age, oxygen level
  - Categorical: contact risk (low/medium/high), fever (yes/no)

### Label / Target (Output Variable)

- The model tries to learn how these features lead to the correct label.

### Training Data

- Portion of the dataset used to teach the model
- The model learns patterns (rules) from this data

### Testing Data

- Portion of the dataset kept aside to check the model's performance
- Ensures the model can handle new/unseen data

### Model

- The mathematical/algorithmic structure that learns patterns from data
- Example: Decision Tree, k-NN, SVM, Logistic Regression

### Prediction

- The model's guess about the target for a new data point

### Performance metrics

- How often the model's predictions are correct

### Overfitting

- When the model memorizes training data instead of learning patterns
- Works perfectly on training data but fails on new data

### Underfitting

Underfitting occurs when a machine learning model is too simple to learn the underlying patterns in the training data. As a result, it performs poorly on both the training data and new (testing) data.

# ARTIFICIAL INTELLIGENCE

---

## ML Supervised Learning

**Dr. Rajani S**

Department of Computer Applications


### ML

- Artificial Intelligence (AI) is a broad field of computer science that focuses on creating machines capable of performing tasks that normally require human intelligence, such as reasoning, learning, problem-solving, perception, and decision-making.
- However, building an AI system by manually programming rules (hard coded) for every possible situation is difficult and often impossible.
- As the amount of data and complexity of problems increased, researchers needed a way for computers to learn automatically from experience instead of being explicitly programmed.
- **This led to the development of Machine Learning (ML)**

### Definition of Machine Learning

Machine Learning is a branch (subset) of Artificial Intelligence that enables computers to learn from data and improve their performance without being explicitly programmed.

Artificial Intelligence aims to make machines intelligent, while Machine Learning is one of the methods used to achieve AI by enabling computers to learn from data rather than relying solely on predefined rules.



## Real-Life Examples

### Artificial Intelligence

---

Self-driving cars

---

Voice assistants

---

Email assistant

---

Netflix recommendations

---

Medical diagnosis

### Machine Learning

---

Learning road signs from images

---

Speech recognition models

---

Spam email detection

---

Predicting movies you may like

---

Disease prediction from patient data

**Machine Learning is needed because it enables systems to:**

**Learn from Data** – Automatically discover patterns instead of relying on manually written rules.

**Adapt Automatically** – Improve performance as more data becomes available.

**Improve with Experience** – Become more accurate over time through training.

**Handle Complex Problems** – Solve tasks where writing explicit rules is difficult or impossible.

**Generalize to New Data** – Make correct predictions for previously unseen examples.


# Types of Machine Learning (ML)

1. Supervised Learning
2. Unsupervised Learning
3. Reinforcement Learning

### Supervised Learning

- Supervised Learning is a task-driven learning approach in which the machine learns from labeled data.
- Each training example contains:
- Input (Features)
- Correct Output (Label)
- The goal is to learn the relationship between the input and output so that the model can predict the correct output for new, unseen data.

SUPERVISED LEARNING



the model **learns** from the **label**

Labelled Data

Cat



Dog



TRAINING



INPUT  
(Test data)



OUTPUT



1. **Input labeled Image:** Cat Image, Dog Image

2. **The machine studies thousands of labeled animal images and learns their characteristics:**

Cats → Pointed ears, whiskers

Dogs → Different face and body structure

Rabbits → Long ears

Birds → Beak and wings

3. **Prediction:** Now the machine receives a new image that it has never seen before.

4. Prediction: **Cat**

Even though the image is new, the model correctly classifies it because it learned from the labeled training data.



Cat




Dog



Bunny

Labelled  
pictures




We are **helping machine** to learn by **labeling each picture** with the name of the animal.

## Unsupervised Learning

- Unsupervised Learning is a type of Machine Learning in which the algorithm learns from unlabeled data.
- Unlike supervised learning, there are no predefined output labels. The algorithm identifies hidden patterns, structures, or relationships within the data on its own.
- **Working of Unsupervised Learning**
- Collect unlabeled data.
- Preprocess the data.
- Apply an unsupervised learning algorithm.
- The algorithm groups similar data or finds hidden patterns.
- Analyze the results.


# ARTIFICIAL INTELLIGENCE


## Machine Learning



### Reinforcement Learning


- Reinforcement Learning (RL) is a type of Machine Learning in which an agent learns by interacting with its environment. The agent performs actions and receives rewards for correct actions or penalties (punishments) for incorrect actions.
- Over time, the agent learns the best strategy (policy) to maximize the total reward.






**1** Observe

**2** Select action using policy



**3** Action!

**4** Get reward or penalty



**5** Update policy (learning step)

**6** Iterate until an optimal policy is found

Agent Plays Game



Makes a Move



Wins? —————▶ Yes → Reward (+100)



No → Penalty (-50)



Learns Better Strategy


### Examples:

Chess AI

Self-driving car simulation

Robot navigation

### Machine Learning Process



### ML workflow or Objectives

- **Collect Data** – Gather relevant raw data from sources
- **Preprocess Data** – Clean, transform, and prepare data for modeling
- **Split Data** – Divide into training and testing sets
- **Choose Model** – Select an appropriate algorithm for the problem
- **Train Model** – Fit the model to the training data
- **Evaluate Model** – Test performance using metrics like accuracy, RMSE,...
- **Deploy Model** – Use the trained model for real-world predictions

**Machine Learning follows a systematic process to build predictive models.**

### Steps

1. **Data Collection** :Data Collection is the process of gathering relevant data from different sources for training the machine learning model.

Manual Observation ,

Fruit Shops / Supermarkets

Farm

Agricultural Research Centers

Public Datasets ( Websites, Databases )

Sensors and Cameras

Collect information about different fruits.

| Weight (g) | Color  | Diameter (cm) | Fruit |
|------------|--------|---------------|-------|
| 160        | Red    | 8             | Apple |
| 180        | Yellow | 9             | Mango |
| 170        | Red    | 8             | Apple |
| 220        | Yellow | 10            | Mango |

**Machine Learning follows a systematic process to build predictive models.**

### Steps

1. **2. Data Preprocessing:** is the process of cleaning and preparing the collected raw data before training.

2. Remove duplicate records

3. Handle missing values

4. Removing noisy or irrelevant data

5. Correct incorrect values

6. Convert categorical data into numerical values (Red =0, green=1)

3: **Feature Selection :** Feature Selection is the process of choosing the most important input features for prediction.

Weight, Color, Diameter

**Machine Learning follows a systematic process to build predictive models.**

### Steps

**4. Model Selection:** Model Selection is the process of choosing a suitable machine learning algorithm for the problem.

1. Decision Tree, k-NN, SVM

**5. Model Training :** During training, the selected algorithm learns patterns from the training dataset.

**6: Model Testing:** The trained model is tested using unseen data to check how well it performs.

Testing Data :

|   | Weight | Color  | Diameter | Actual Fruit |
|---|--------|--------|----------|--------------|
|   |        |        |          |              |
|   | 165    | Red    | 8        | Apple        |
| 1 | 205    | Yellow | 9        | Mango        |

**Machine Learning follows a systematic process to build predictive models.**

## Steps

### 7. Prediction:

| Actual | Predicted |
|--------|-----------|
| Apple  | Apple     |
| Mango  | Mango     |

1.

### 8. Model Evaluation:

Model Evaluation measures the performance of the trained model using evaluation metrics.

### 9. Deployment :

The fruit shop can now automatically classify new fruits.

**Machine Learning follows a systematic process to build predictive models.**

### **Disease Prediction:**

Collect patient records.

Remove missing values., Handle duplicate records , Normalize data

Select features like age, BP, sugar.

Train Decision Tree.

Test using unseen data.

Deploy in hospitals.

Machine Learning is revolutionizing many sectors.

| Industry       | Application                        |
|----------------|------------------------------------|
| Healthcare     | Disease diagnosis, medical imaging |
| Banking        | Fraud detection, loan approval     |
| Retail         | Product recommendation             |
| Agriculture    | Crop disease prediction            |
| Manufacturing  | Predictive maintenance             |
| Transportation | Self-driving cars                  |
| Education      | Personalized learning              |
| Entertainment  | Movie recommendations              |



**THANK YOU**

---

**Dr. Rajani S**

Department of Computer Applications

[rajanis@pes.edu](mailto:rajanis@pes.edu)